

Using the Dictionary Structure for efficient Convolutional Dictionary Learning

Thomas Moreau – INRIA – Université Paris Saclay



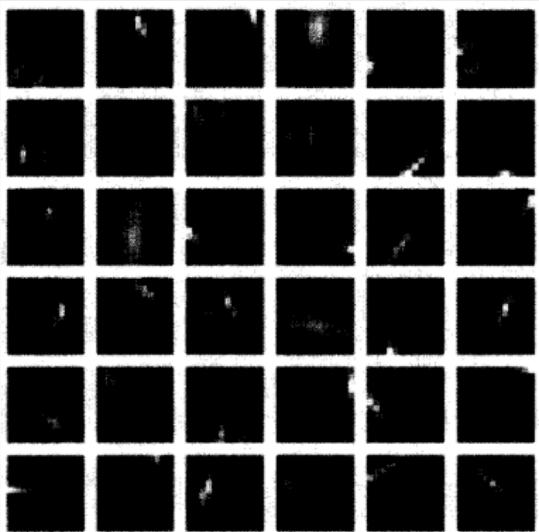
PARIETAL

informatics mathematics
inria

Dictionary learning learns a set of atoms (patterns) to sparsely reconstruct a signal,


Goal:

- ▶ Feature extraction,
- ▶ Signal exploration.

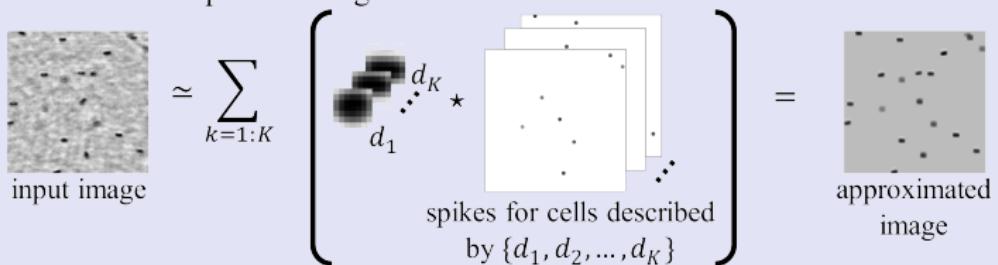


Patches learned with natural images in Olshausen and Field 1997.

Sparse coding



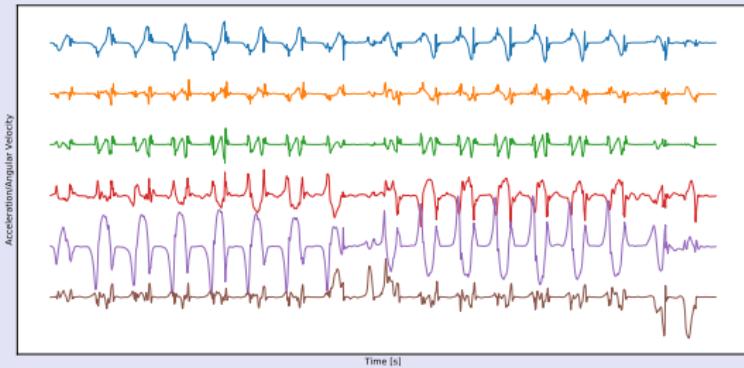
Convolutional sparse coding



Convolutional Dictionary Learning learns a set of **shift-invariant** atoms to sparsely reconstruct a signal,

- ▶ Improve sparsity
- ▶ Not all patches are encoded
- ▶ Sharper atoms

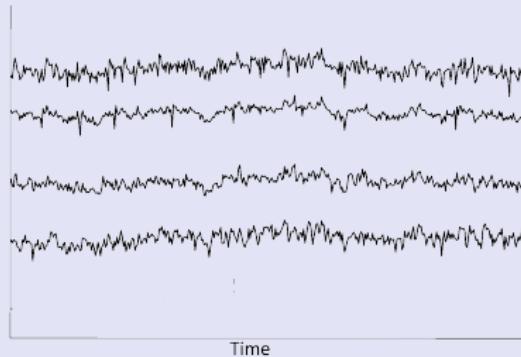
Application fields



[Oudre et al. 2018, Sensors]

- ▶ Detecting steps in human walk recordings to predict elderly falls.

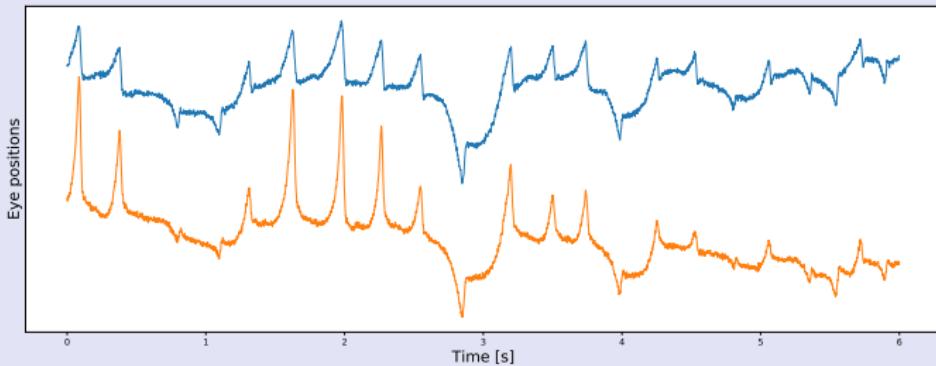
Application fields



[Dupré la Tour et al. 2018, NeurIPS]

- ▶ Detecting steps in human walk recordings to predict elderly falls.
- ▶ Exploring neurological signals from ECG and MEG,

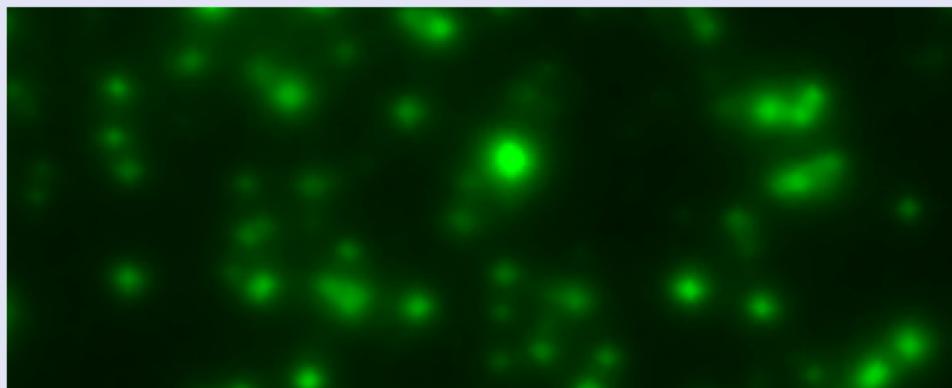
Application fields



[Robert et al. 2016, preprint]

- ▶ Detecting steps in human walk recordings to predict elderly falls.
- ▶ Exploring neurological signals from ECG and MEG,
- ▶ Classifying pathological eye movements from oculomotor signals.

Application fields



[del Aguila Pla et al. 2018, IEEE TSP; Yellin et al. 2017, ISBI]

- ▶ Detecting steps in human walk recordings to predict elderly falls.
- ▶ Exploring neurological signals from ECG and MEG,
- ▶ Classifying pathological eye movements from oculomotor signals.
- ▶ Counting cells in biological images.

Application fields



[del Aguila Pla et al. 2018, ICASSP;
Beckouche et al. 2013, Astronomy & Astrophysics]

- ▶ Detecting steps in human walk recordings to predict elderly falls.
- ▶ Exploring neurological signals from ECG and MEG,
- ▶ Classifying pathological eye movements from oculomotor signals.
- ▶ Counting cells in biological images.
- ▶ Counting stars and galaxies in telescope images

Challenges of Convolutional Dictionary Learning

- ▶ **Computational:** how to scale with large signals,
 - ▶ by exploiting the structure of the dictionary.
 - ▶ by parallelization.
- ▶ **Modelization:** how to incorporate prior knowledge,
 - ▶ on the activations.
 - ▶ on the patterns.
- ▶ **Evaluation:** how to evaluate the quality of the learned patterns.
- ▶ **Theoretical:** pattern recovery.

Challenges of Convolutional Dictionary Learning

- ▶ **Computational:** how to scale with large signals,
 - ▶ by exploiting the structure of the dictionary.
[Moreau and Bruna 2017, ICLR]
 - ▶ by parallelization.
[Moreau et al. 2018, ICML; Moreau and Gramfort 2019, preprint]
- ▶ **Modelization:** how to incorporate prior knowledge,
 - ▶ on the activations.
 - ▶ on the patterns.
[Dupré la Tour et al. 2018, NeurIPS]
- ▶ **Evaluation:** how to evaluate the quality of the learned patterns.
- ▶ **Theoretical:** pattern recovery.

Convolutional Dictionary Learning

Adaptive Sparse Coding

Scaling up Convolutional Sparse Coding with
coordinate descent and distributed optimization

Rank-1 Constrained Convolutional Dictionary Learning

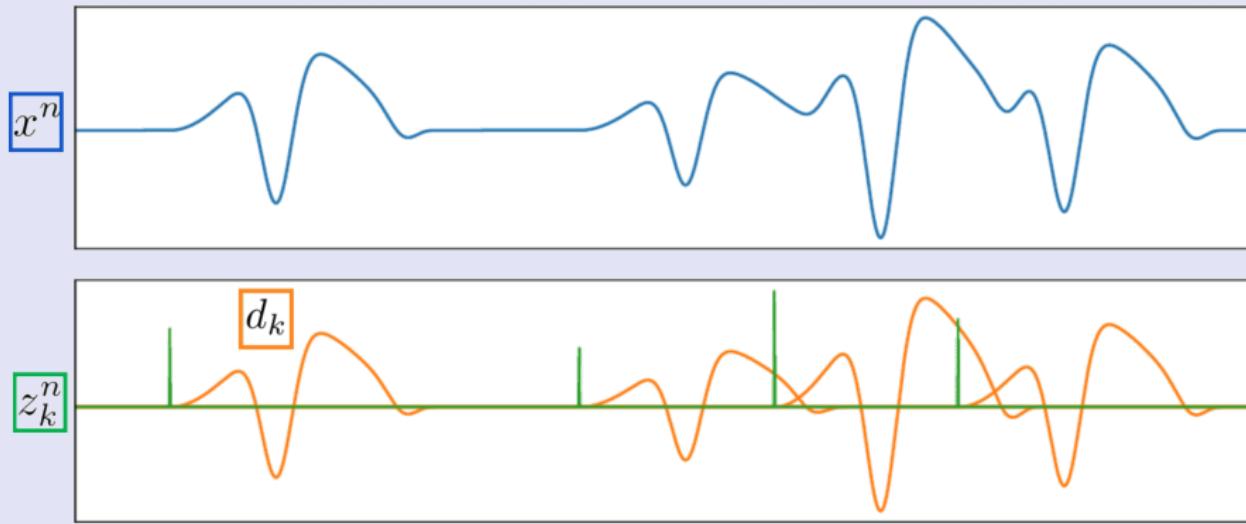
Convolutional Dictionary Learning

References

- ▶ Grosse, R., Raina, R., Kwong, H., and Ng, A. Y. (2007). Shift-Invariant Sparse Coding for Audio Classification. *Cortex*, 8:9

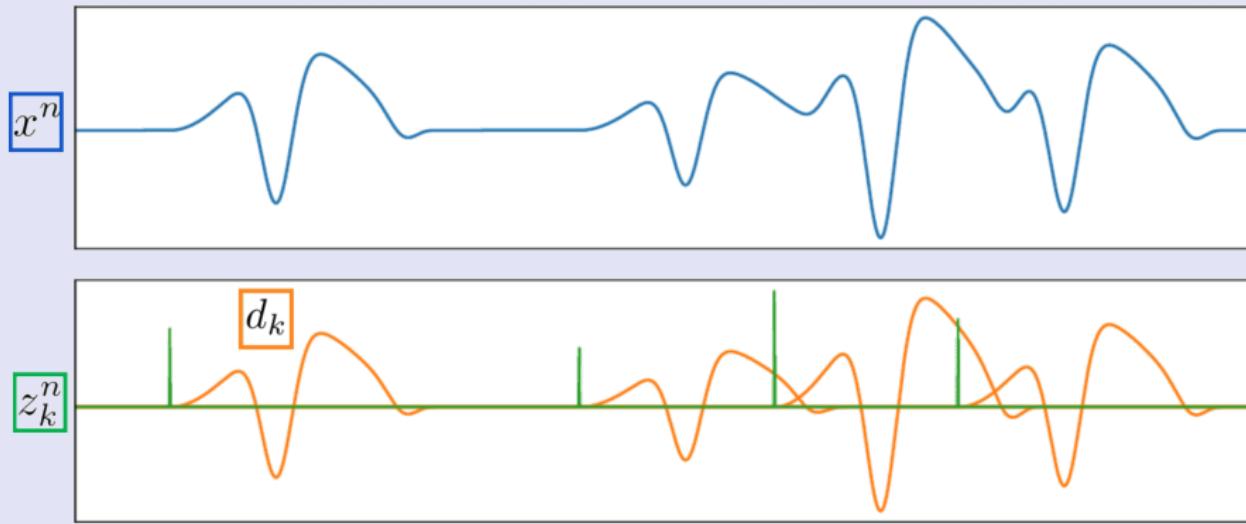
Extracting shift invariant patterns

Key idea: decouple the localization of the patterns and their shape



Extracting shift invariant patterns

Key idea: decouple the localization of the patterns and their shape

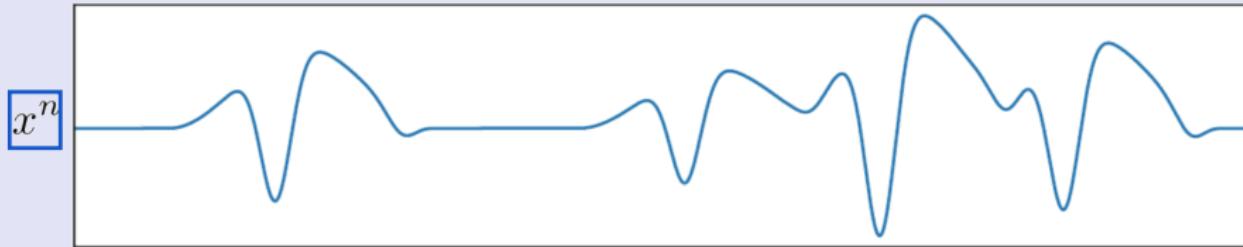


**Convolutional
Representation:**

$$x^n[t] = \sum_{k=1}^K (z_k^n * d_k)[t] + \varepsilon[t]$$

Extracting shift invariant patterns

Key idea: decouple the localization of the patterns and their shape



**Convolutional
Dictionary Learning:**

$$\begin{aligned} & \min_{d,z} \sum_{n=1}^N \frac{1}{2} \left\| x^n - \sum_{k=1}^K z_k^n * d_k \right\|_2^2 + \lambda \sum_{k=1}^K \| z_k^n \|_1, \\ & \text{s.t. } \| d_k \|_2^2 \leq 1 \end{aligned}$$

Notation

Sparse Convolutional model:

$$X[t] = \sum_{k=1}^K (\mathcal{D}_k * Z_k)[t] + \mathcal{E}[t]$$

with Z sparse. Few of its coefficients are non-zero.

- ▶ X is a signal of length T
- ▶ \mathcal{E} is a noise signal of length T
- ▶ \mathcal{D} is a set of K patterns of length L
- ▶ Z is a signal of length $\tilde{T} = T - L + 1$ in \mathbb{R}^K

Convolutional Dictionary Learning

Dictionary learning optimization problem for $\{X^{[n]}\}_{n=1}^N$

$$\min_{Z, \|\mathcal{D}_k\| \leq 1} \frac{1}{N} \sum_{n=1}^N \underbrace{\left\| X^{[n]} - \sum_{k=1}^K \mathcal{D}_k * Z_k^{[n]} \right\|_2^2}_{E(Z) \text{ data fit}} + \underbrace{\lambda \|Z^{[n]}\|_1}_{\text{penalization}}$$

with a regularization parameter $\lambda > 0$.

This problem is bi-convex and an approximate solution is obtained through
alternate minimization. [Engan et al., 1999; Grosse et al., 2007]

D-step: Dictionary updates

→ Z fixed, update \mathbf{D}

$$\mathbf{D}^* = \operatorname*{argmin}_{\|\mathbf{D}_k\|_2 \leq 1} \frac{1}{N} \sum_{n=1}^N \|X^{[n]} - \sum_{k=1}^K \mathbf{D}_k * Z_k^{[n]}\|_2^2$$

Related Algorithms:

- ▶ Proximal Gradient Descent (PDG) [Rockafellar, 1976]
- ▶ Accelerated PGD [Nesterov, 1983]
- ▶ Block Coordinate Descent [Mairal et al., 2010]
- ▶ Alternated Direction Method of Multiplier (ADMM)
[Gabay and Mercier, 1976]

Z -step: Convolutional Sparse Coding (CSC)

$\rightarrow \mathbf{D}$ fixed, update Z

$$Z^{[n],*} = \operatorname{argmin}_{Z^{[n]}} \|X^{[n]} - \sum_{k=1}^K \mathbf{D}_k * Z_k^{[n]}\|_2^2 + \lambda \|Z^{[n]}\|_1$$

\Rightarrow Independent for each $n \in \llbracket 1, N \rrbracket$

Related Algorithms:

- ▶ Iterative Soft-Thresholding Algorithm (ISTA)
[Daubechies et al., 2004; Chalasani et al., 2013]
- ▶ Fast ISTA
[Beck and Teboulle, 2009; Wohlberg, 2016]
- ▶ Alternated Direction Method of Multiplier (ADMM)
[Gabay and Mercier, 1976; Bristow et al., 2013]
- ▶ Coordinate Descent (CD)
[Friedman et al., 2007; Kavukcuoglu et al., 2010]

Adaptive Sparse Coding

References

- Moreau, T. and Bruna, J. (2017). Understanding Neural Sparse Coding with Matrix Factorization. In *International Conference on Learning Representation (ICLR)*

Adaptive Optimization for the Z -step

We have to solve N problems with a common structure D .

$$Z^{[n],*} = \operatorname{argmin}_{Z^{[n]}} \|X^{[n]} - \sum_{k=1}^K D_k * Z_k^{[n]}\|_2^2 + \lambda \|Z^{[n]}\|_1$$

Can we use this structure to accelerate the resolution?

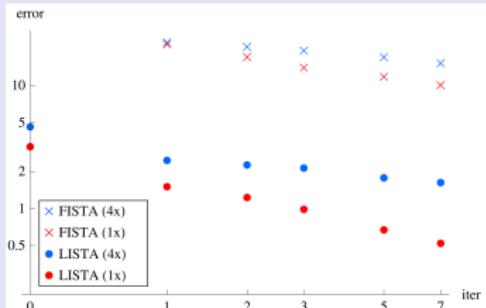
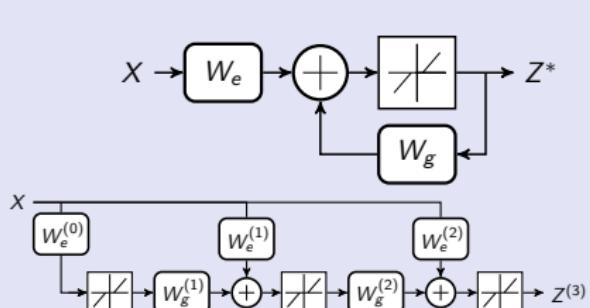
Adaptive Optimization for the Z -step

We have to solve N problems with a common structure D .

$$Z^{[n],*} = \underset{Z^{[n]}}{\operatorname{argmin}} \|X^{[n]} - \sum_{k=1}^K D_k * Z_k^{[n]}\|_2^2 + \lambda \|Z^{[n]}\|_1$$

Can we use this structure to accelerate the resolution?

Yes, with the Learned ISTA [Gregor and Le Cun 2010, NeurIPS]



Adaptive Optimization for the Z -step

We have to solve N problems with a common structure D .

$$Z^{[n],*} = \operatorname{argmin}_{Z^{[n]}} \|X^{[n]} - \sum_{k=1}^K D_k * Z_k^{[n]}\|_2^2 + \lambda \|Z^{[n]}\|_1$$

Can we use this structure to accelerate the resolution?

Yes, with the Learned ISTA [Gregor and Le Cun 2010, NeurIPS]

Why does it work? (non-convolutional setting).

- ▶ [Xin et al. 2016, NeurIPS]: Improved support recovery.
- ▶ [Giryes et al. 2018, IEEE TSP]: Leverage the sparsity pattern in Z .
- ▶ [Chen et al. 2018, NeurIPS]: Linear convergence if Z sparse enough.

Adaptive Optimization for the Z -step

We have to solve N problems with a common structure D .

$$Z^{[n],*} = \operatorname{argmin}_{Z^{[n]}} \|X^{[n]} - \sum_{k=1}^K D_k * Z_k^{[n]}\|_2^2 + \lambda \|Z^{[n]}\|_1$$

Can we use this structure to accelerate the resolution?

Yes, with the Learned ISTA [Gregor and Le Cun 2010, NeurIPS]

Why does it work? (non-convolutional setting).

- ▶ [Xin et al. 2016, NeurIPS]: Improved support recovery.
- ▶ [Giryes et al. 2018, IEEE TSP]: Leverage the sparsity pattern in Z .
- ▶ [Chen et al. 2018, NeurIPS]: Linear convergence if Z sparse enough.
- ▶ [Moreau and Bruna 2017, ICLR]: Leverage the structure of D .

Vectorized model

- ▶ x is a vector in \mathbb{R}^T
- ▶ ϵ is a noise vector in \mathbb{R}^T
- ▶ D is a matrix in $\mathbb{R}^{T \times LK}$
- ▶ z is a coding vector in $\mathbb{R}^{\tilde{T}K}$

Sparse Linear model:

$$x = Dz + \epsilon$$

with z sparse.

Few of its coefficients are non-zero.

$$D = \left(\begin{array}{c|c} \begin{matrix} \bar{D} & 0 \\ 0 & \bar{D} \\ 0 & 0 \\ 0 & \bar{D} \end{matrix} & 0 \\ \hline 0 & \begin{matrix} 0 & \bar{D} & 0 \\ 0 & \bar{D} & 0 \\ 0 & 0 & \bar{D} \end{matrix} \end{array} \right)$$

Notations

Consider the sparse coding problem with a dictionary D .

$$z^* = \operatorname{argmin}_z F(z) = \underbrace{\frac{1}{2} \|x - Dz\|_2^2}_{E(z)} + \lambda \|z\|_1$$

We denote $B = D^T D$ is the Gram matrix of D .

We introduce a novel class of algorithms – FacNet – based on a sparse factorization of B .

Quadratic form: $Q_S(u, v) = \frac{1}{2}(u - v)^T S(u - v) + \lambda \|u\|_1$.

Note that $F(z) = Q_B(z, D^\dagger x)$ and $\operatorname{prox}_{\|\cdot\|_1}^S(v) = \operatorname{argmin}_u Q_S(u, v)$

If S is diagonal, $\operatorname{argmin}_u Q_S(u, v)$ can be efficiently minimized as the problem is separable in each coordinate.

Given an estimate $z^{(q)}$ of z^* at iteration q , we can write:

$$\begin{aligned} F(z) &= E(z) + \lambda \|z\|_1 \\ &= E(z^{(q)}) + \left\langle \nabla E(z^{(q)}), z - z^{(q)} \right\rangle + Q_B(z, z^{(q)}) , \end{aligned}$$

Given an estimate $z^{(q)}$ of z^* at iteration q , we can write:

$$\begin{aligned} F(z) &= E(z) + \lambda \|z\|_1 \\ &= E(z^{(q)}) + \left\langle \nabla E(z^{(q)}), z - z^{(q)} \right\rangle + Q_B(z, z^{(q)}) , \end{aligned}$$

ISTA: Replace B by diagonal matrix $\textcolor{red}{S} = \|B\|_2 \mathbf{I}_K$

$$F_q(z) = E(z^{(q)}) + \left\langle \nabla E(z^{(q)}), z - z^{(q)} \right\rangle + Q_{\textcolor{red}{S}}(z, z^{(q)}) ,$$

$$\min_z F_q(z) \Leftrightarrow \min_z Q_{\textcolor{red}{S}} \left(z, z^{(q)} - \textcolor{red}{S}^{-1} \nabla E(z^{(q)}) \right)$$

Given an estimate $z^{(q)}$ of z^* at iteration q , we can write:

$$\begin{aligned} F(z) &= E(z) + \lambda \|z\|_1 \\ &= E(z^{(q)}) + \left\langle \nabla E(z^{(q)}), z - z^{(q)} \right\rangle + Q_B(z, z^{(q)}) , \end{aligned}$$

ISTA: Replace B by diagonal matrix $\textcolor{red}{S} = \|B\|_2 \mathbf{I}_K$

FacNet: Replace B by $\textcolor{blue}{A}_q^\top \textcolor{blue}{S}_q \textcolor{blue}{A}_q$ ($\textcolor{blue}{S}_q$ diagonal, $\textcolor{blue}{A}_q$ unitary)

$$\widetilde{F}_q(z) = E(z^{(q)}) + \left\langle \nabla E(z^{(q)}), z - z^{(q)} \right\rangle + Q_{\textcolor{blue}{S}_q}(\textcolor{blue}{A}_q z, \textcolor{blue}{A}_q z^{(q)}) ,$$

$$\min_z \widetilde{F}_q(z) \Leftrightarrow \min_z Q_{\textcolor{blue}{S}_q} \left(\textcolor{blue}{A}_q z, \textcolor{blue}{A}_q z^{(q)} - \textcolor{blue}{S}_q^{-1} \textcolor{blue}{A}_q \nabla E(z^{(q)}) \right)$$

Given an estimate $z^{(q)}$ of z^* at iteration q , we can write:

$$\begin{aligned} F(z) &= E(z) + \lambda \|z\|_1 \\ &= E(z^{(q)}) + \left\langle \nabla E(z^{(q)}), z - z^{(q)} \right\rangle + Q_B(z, z^{(q)}) , \end{aligned}$$

ISTA: Replace B by diagonal matrix $\textcolor{red}{S} = \|B\|_2 \mathbf{I}_K$

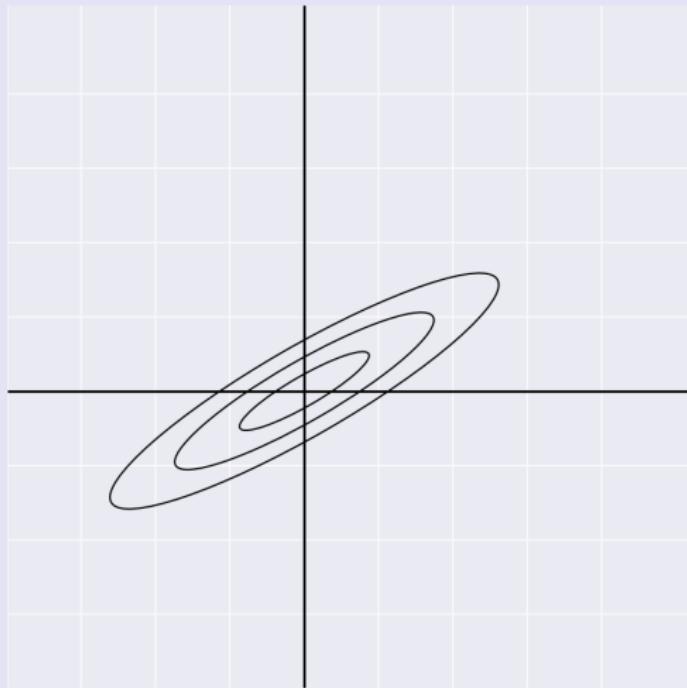
FacNet: Replace B by $\textcolor{blue}{A}_q^\top \textcolor{blue}{S}_q \textcolor{blue}{A}_q$ ($\textcolor{blue}{S}_q$ diagonal, $\textcolor{blue}{A}_q$ unitary)

$$\widetilde{F}_q(z) = E(z^{(q)}) + \left\langle \nabla E(z^{(q)}), z - z^{(q)} \right\rangle + Q_{\textcolor{blue}{S}_q}(\textcolor{blue}{A}_q z, \textcolor{blue}{A}_q z^{(q)}) ,$$


$$\min_z \widetilde{F}_q(z) \Leftrightarrow \min_z Q_{\textcolor{blue}{S}_q} \left(\textcolor{blue}{A}_q z, \textcolor{blue}{A}_q z^{(q)} - \textcolor{blue}{S}_q^{-1} \textcolor{blue}{A}_q \nabla E(z^{(q)}) \right)$$

Can we choose A_q, S_q to accelerate the optimization compared to ISTA?

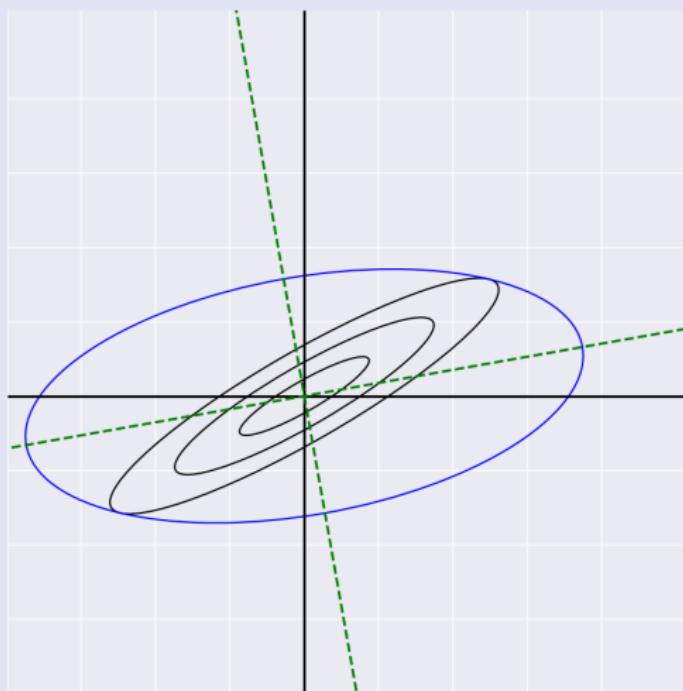
Toward an adaptive procedure



Toward and adaptive procedure



Toward and adaptive procedure



Toward an adaptive procedure

The surrogate \widetilde{F}_q can be re-written as

$$\widetilde{F}_q(z) = F(z) + (z - z^{(q)})^T R(z - z^{(q)}) + \delta_A(z) .$$

Tradeoff between:

- ▶ Diagonalization of the gram matrix B , Computation

$$R = A^T S A - B$$

- ▶ Deformation of the ℓ_1 -norm with the rotation A . Accuracy

$$\delta_A(z) = \lambda \left(\|Az\|_1 - \|z\|_1 \right)$$

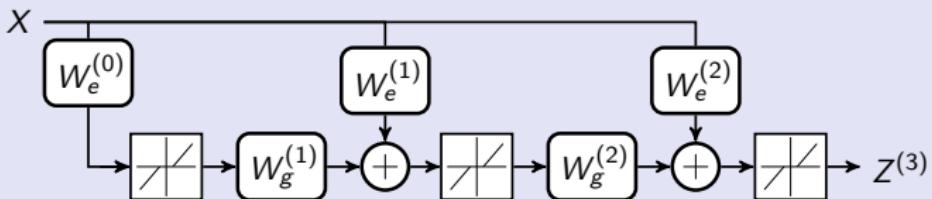
Theoretical results

- ▶ We showed that FacNet has the same asymptotic convergence rate as ISTA in $\mathcal{O}(\frac{1}{q})$.
- ▶ The constant factors are different and can be improved. If the factorization (A_q, S_q) at iteration q verifies

$$\|R_q\|_2 + 2 \frac{L_{A_q}(z^{(q+1)})}{\|z^* - z^{(q)}\|_2} \leq \frac{\|B\|_2}{2}$$

and $A_p = \mathbf{I}_K, S_p = \|B\|_2 \mathbf{I}_K$ for $p > q$, then the procedure has improved convergence rate compared to ISTA.

⇒ There is a phase transition when $\|z^{(q)} - z^*\|_2 \rightarrow 0$

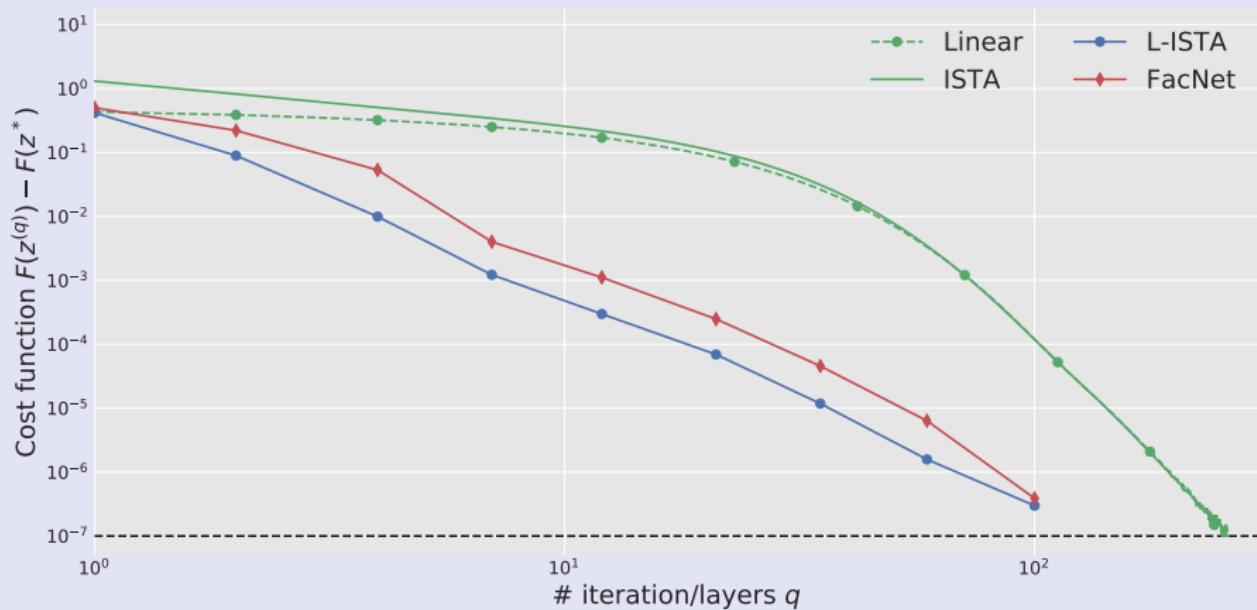


With $W_e = \frac{D^T}{\|B\|_2}$ and $W_g = I - \frac{B}{\|B\|_2}$, this network computes ISTA.

FacNet: Specialization of LISTA with $\begin{cases} W_e &= S^{-1}AD^T \\ W_g &= A^T - S^{-1}ABA^T \end{cases}$

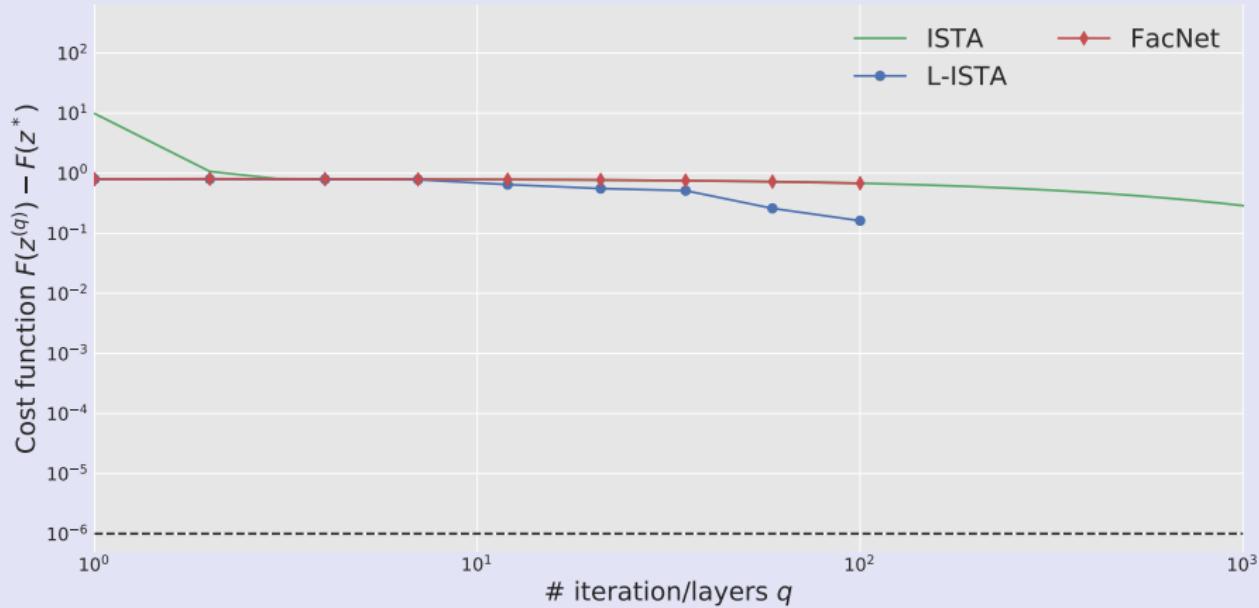
⇒ LISTA can be at least as good as this model.

Generic Dictionary



K generic atoms (uniform in S^{P-1}) with Bernoulli-Gaussian activation.
Params: $K = 100$, $P = 64$, $\rho = 1/20$, $\sigma = 10$ and $\lambda = 0.01$

Adversarial dictionary



Same parameters with adverse dictionary (dense eigen-spaces).

Take home message

- ▶ Non asymptotic acceleration of ISTA is possible based on the structure of D ,
- ▶ Sufficient analysis to explain LISTA acceleration,
- ▶ Empirically showed the structure of D is necessary for LISTA.

Ahead of us

- ▶ Improve the factorization formulation for direct optimization,
- ▶ Adaptation of the analysis to convolutional sparse coding,
- ▶ Explore the link with sparse eigenvectors of the gram matrix.

Scaling up Convolutional Sparse Coding with coordinate descent and distributed optimization

References

- ▶ Moreau, T., Oudre, L., and Vayatis, N. (2018). [DICOD: Distributed Convolutional Sparse Coding](#). In *International Conference on Machine Learning (ICML)*, pages 3626–3634, Stockholm, Sweden. PMLR (80)
- ▶ Moreau, T. and Gramfort, A. (2019). [Distributed Convolutional Dictionary Learning \(DiCoDiLe\): Pattern Discovery in Large Images and Signals](#). *preprint ArXiv*, 1901.09235

Z -step: Sparse coding

$\rightarrow \mathbf{D}$ fixed, update Z

$$Z^{[n],*} = \operatorname*{argmin}_{Z^{[n]}} \|X^{[n]} - \sum_{k=1}^K \mathbf{D}_k * Z_k^{[n]}\|_2^2 + \lambda \|Z^{[n]}\|_1$$

\Rightarrow Independent for each $n \in [1, N]$

Related Algorithms:

- ▶ Iterative Soft-Thresholding Algorithm (ISTA)
[Daubechies et al., 2004; Chalasani et al., 2013]
- ▶ Fast ISTA
[Beck and Teboulle, 2009; Wohlberg, 2016]
- ▶ Alternated Direction Method of Multiplier (ADMM)
[Gabay and Mercier, 1976; Bristow et al., 2013]
- ▶ Coordinate Descent (CD)
[Friedman et al., 2007; Kavukcuoglu et al., 2010]

Coordinate Descent (CD)

Minimize

$$Z^* = \operatorname{argmin}_Z \|X - \sum_{k=1}^K D_k * Z_k\|_2^2 + \lambda \|Z\|_1$$

Update one coordinate at each iteration.

1. Select a coordinate (k_0, t_0) to update.

Three algorithms for LASSO:

- ▶ Cyclic updates; $\mathcal{O}(1)$ [Friedman et al., 2007]
- ▶ Random updates; $\mathcal{O}(1)$ [Nesterov, 2010]
- ▶ Greedy updates; $\mathcal{O}(KL)$ [Osher and Li, 2009]

Coordinate Descent (CD)

Minimize

$$Z^* = \operatorname{argmin}_Z \|X - \sum_{k=1}^K D_k * Z_k\|_2^2 + \lambda \|Z\|_1$$

Update one coordinate at each iteration.

1. Select a coordinate (k_0, t_0) to update.
2. Compute a new value $Z'_{k_0}[t_0]$ for this coordinate

For convolutional CD, we can use optimal updates:

$$Z'_{k_0}[t_0] = \frac{1}{\|D_{k_0}\|_2^2} \mathbf{ST}(\beta_{k_0}[t_0], \lambda),$$

with $\mathbf{ST}(y, \lambda) = \operatorname{sign}(y)(|y| - \lambda)_+$. [Kavukcuoglu et al. \[2010\]](#) showed this can be done efficiently, with $\mathcal{O}(KW)$ operations.

Coordinate Descent (CD)

Minimize

$$Z^* = \operatorname{argmin}_Z \|X - \sum_{k=1}^K D_k * Z_k\|_2^2 + \lambda \|Z\|_1$$

Update one coordinate at each iteration.

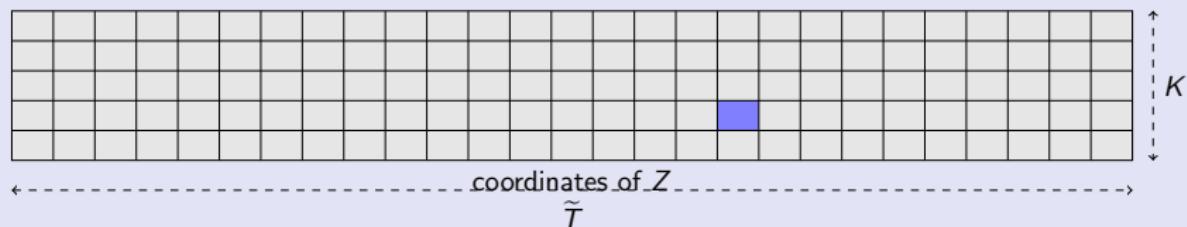
1. Select a coordinate (k_0, t_0) to update.
2. Compute a new value $Z'_{k_0}[t_0]$ for this coordinate

\Rightarrow Converges to the optimal point for CSC problem in $\mathcal{O}\left(\frac{1}{q}\right)$ iterations.

Trade-off between cheap computational complexity (random/cyclic CD) and importance sampling with faster convergence (Greedy CD).

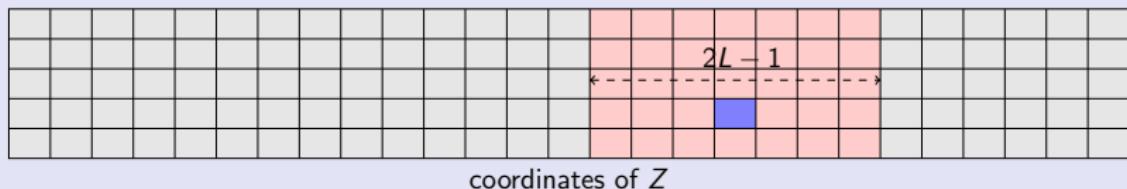
[Nutini et al., 2015; Karimireddy et al., 2019]

We introduced the LGCD method which is an extension of GCD.



GCD has $\mathcal{O}(K\tilde{T})$ computational complexity.

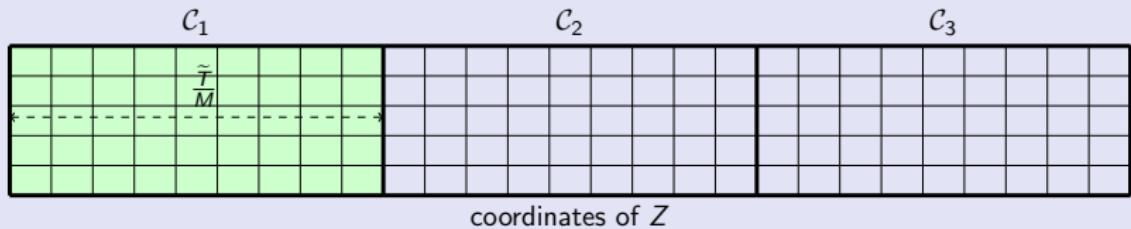
We introduced the LGCD method which is an extension of GCD.



GCD has $\mathcal{O}(K\tilde{T})$ computational complexity.

But the update itself has complexity $\mathcal{O}(KL)$

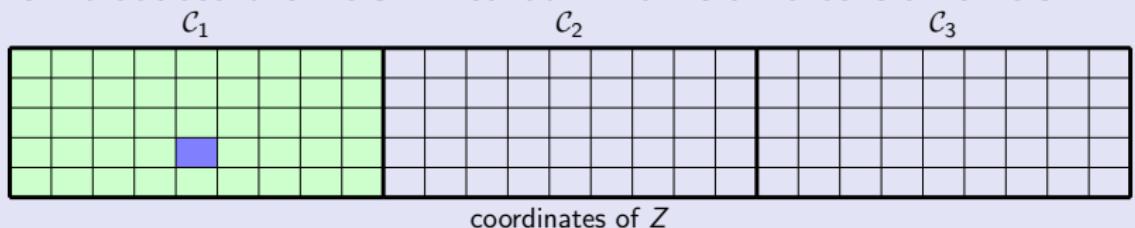
We introduced the LGCD method which is an extension of GCD.



With a partition \mathcal{C}_m of the signal domain $[1, K] \times [0, \tilde{T}]$,

$$\mathcal{C}_m = [1, K] \times \left[\frac{(m-1)\tilde{T}}{M}, \frac{m\tilde{T}}{M} \right]$$

We introduced the LGCD method which is an extension of GCD.



With a partition \mathcal{C}_m of the signal domain $[1, K] \times [0, \tilde{T}]$,

$$\mathcal{C}_m = [1, K] \times \left[\frac{(m-1)\tilde{T}}{M}, \frac{m\tilde{T}}{M} \right]$$

The coordinate to update is chosen greedily on a sub-domain \mathcal{C}_m

$$\frac{\tilde{T}}{M} = 2L - 1 \Rightarrow \mathcal{O}(\text{Coordinate selection}) = \mathcal{O}(\text{Coordinate Update})$$

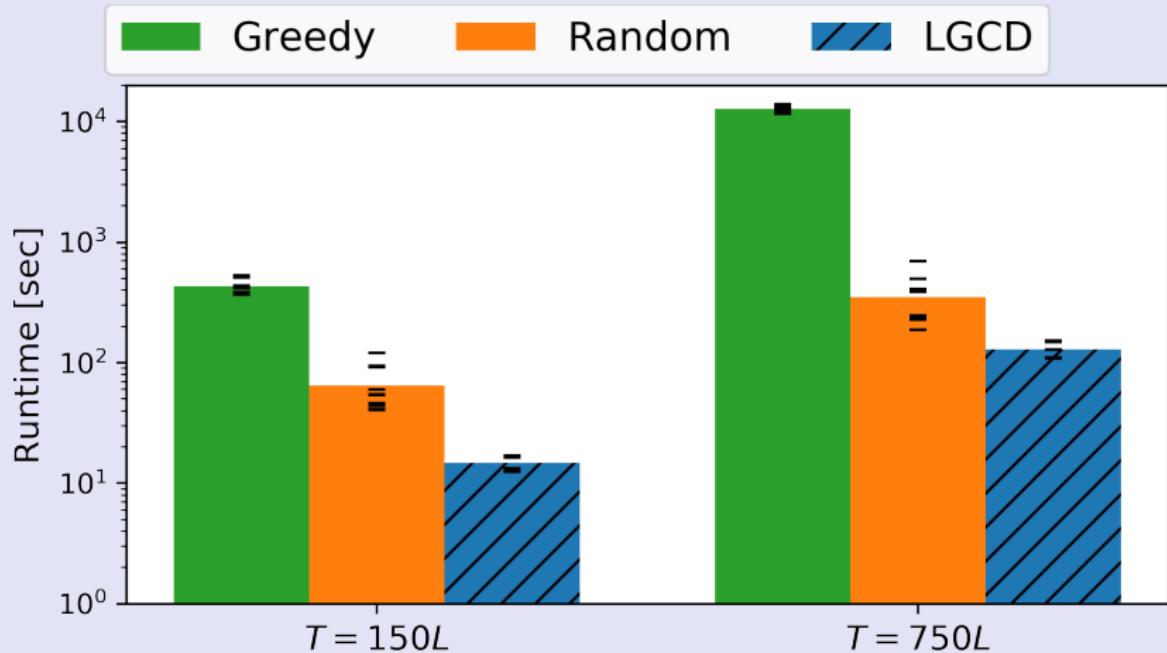
The overall iteration complexity is $\mathcal{O}(KL)$ instead of $\mathcal{O}(K\tilde{T})$.

\Rightarrow Efficient for sparse Z

Fast optimization

Comparison of the coordinate selection strategy for CD on simulated signals

We set $K = 10$, $L = 150$, $\lambda = 0.1\lambda_{\max}$



Weak dependence of the coordinate updates

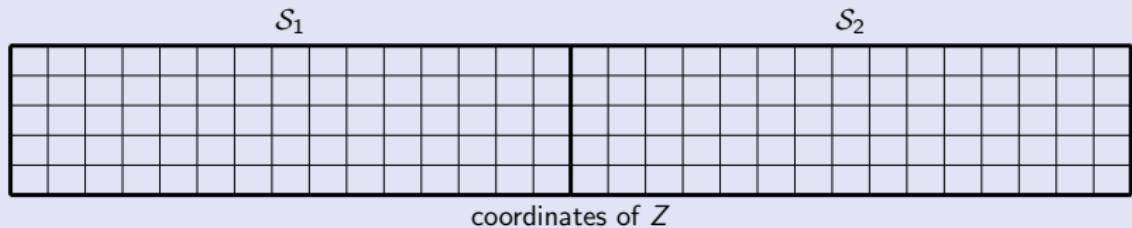
The update of the W coordinates $(k_w, \omega_w)_{w=1}^W$ with additive update $\Delta Z_{k_w}[\omega_w]$ changes the cost by:

$$\Delta E = \underbrace{\sum_{i=1}^W \Delta E_w}_{\text{iterative steps}} - \underbrace{\sum_{w \neq w'} (d_{k_w} * d_{k_{w'}}^\top) [\omega_{w'} - \omega_w] \Delta Z_{k_w}[\omega_w] \Delta Z_{k_{w'}}[\omega_{w'}]}_{\text{interference}},$$

⇒ If the updates are far enough, they can be considered as independent.

Distributed Convolutional Coordinate Descent

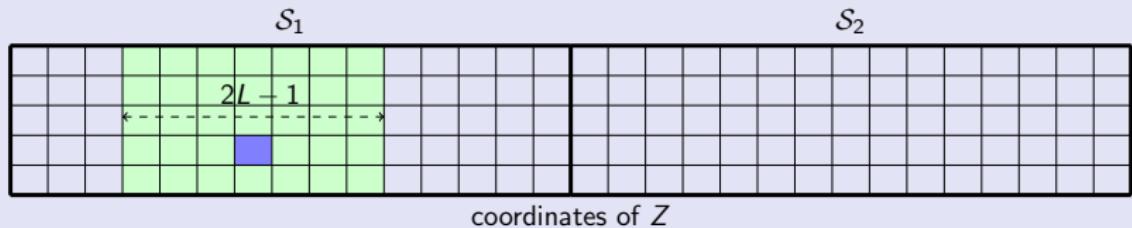
[Moreau et al. 2018, ICML]



- ▶ Split the coordinates in continuous sub-segment $\mathcal{S}_w = \left[\frac{(w-1)T}{W}, \frac{wT}{W} \right]$

Distributed Convolutional Coordinate Descent

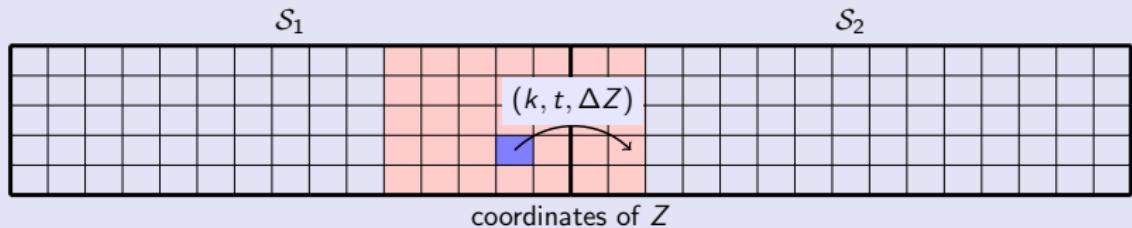
[Moreau et al. 2018, ICML]



- ▶ Split the coordinates in continuous sub-segment $\mathcal{S}_w = \left[\frac{(w-1)T}{W}, \frac{wT}{W} \right]$.
- ▶ Use CD updates in parallel in each sub-segment.

Distributed Convolutional Coordinate Descent

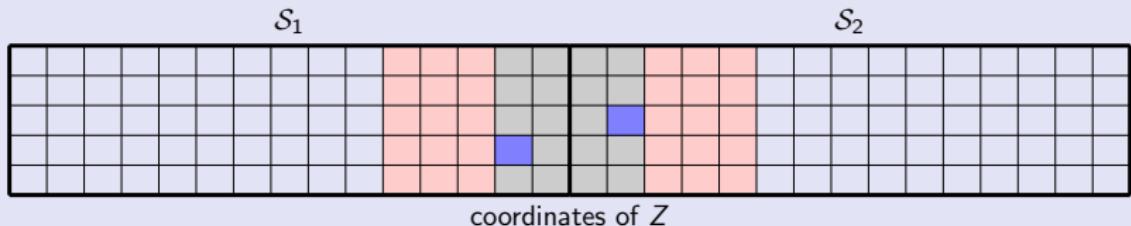
[Moreau et al. 2018, ICML]



- ▶ Split the coordinates in continuous sub-segment $\mathcal{S}_w = \left[\frac{(w-1)T}{W}, \frac{wT}{W} \right]$.
- ▶ Use CD updates in parallel in each sub-segment.
- ▶ Notify neighbor workers when the update is on the border of \mathcal{S}_w .

Distributed Convolutional Coordinate Descent

[Moreau et al. 2018, ICML]



- ▶ Split the coordinates in continuous sub-segment $\mathcal{S}_w = \left[\frac{(w-1)T}{W}, \frac{wT}{W} \right]$.
- ▶ Use CD updates in parallel in each sub-segment.
- ▶ Notify neighbor workers when the update is on the border of \mathcal{S}_w .
- ▶ What do we do when two updates are interfering?

DICOD converges to the solution of the CSC for 1D signals without having a control mechanism on the interference.

Theorem (Convergence of DICOD)

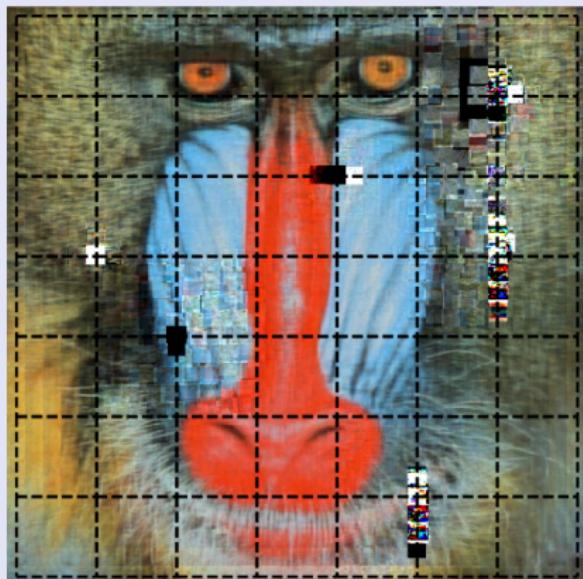
We consider the following assumptions:

- H1:** If the cross correlation between atoms of \mathcal{D} is strictly smaller than 1.
- H2:** No cores stop before all its coefficients are optimal.
- H3:** If the delay in communication between the processes is inferior to the update time.

Under these assumptions, the DICOD algorithm converges asymptotically to the optimal solution Z^* of CSC.

Distributed Convolutional Dictionary Learning (DiCoDiLe-Z)


[Moreau and Gramfort 2019, preprint]



- ▶ DICOD does not work for higher dimensional signals.

Distributed Convolutional Dictionary Learning (DiCoDiLe-Z)

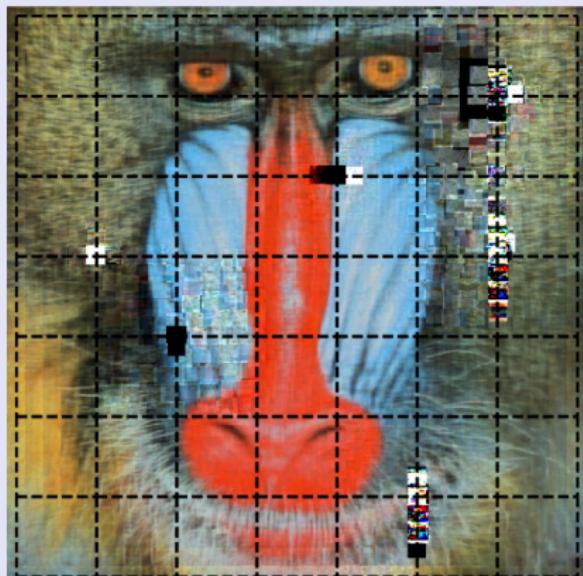
[Moreau and Gramfort 2019, preprint]



- ▶ DICOD does not work for higher dimensional signals.
- ▶ Extension require to control interferences.

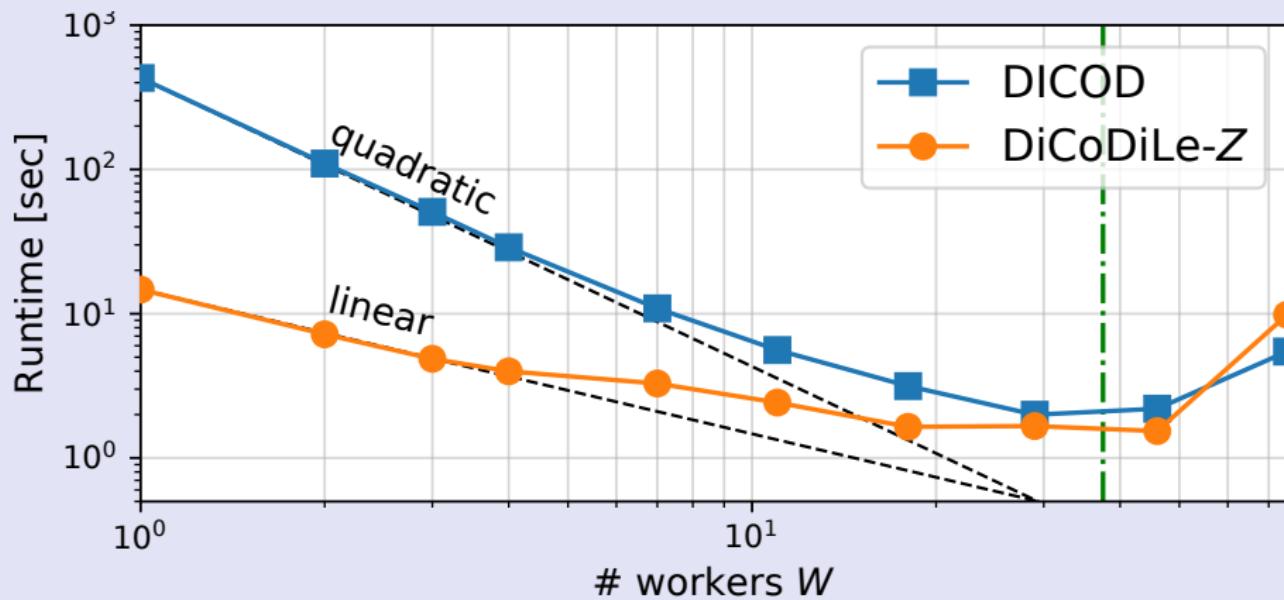
Distributed Convolutional Dictionary Learning (DiCoDiLe-Z)

[Moreau and Gramfort 2019, preprint]



- ▶ DICOD does not work for higher dimensional signals.
- ▶ Extension require to control interferences.
- ▶ Use asynchronous mechanism:
Soft-lock.

Numerical speed-up



Running time as a function fo the number of workers W .

Recap Part II

Take home message

- ▶ LGCD is a very efficient algorithm when working with CSC for long signals.
- ▶ Can be distributed efficiently for multi-dimensional signals,
- ▶ Good scaling properties with the number of workers W used to distribute the algorithm.

Ahead of us

- ▶ Extend this algorithm to local penalization such as Group LASSO.
- ▶ This algorithm could be used for algorithm such as MP for ℓ_0 or $\ell_{0,\infty}$ penalties.

Rank-1 Constrained Convolutional Dictionary Learning

References

- ▶ Dupré la Tour, T., Moreau, T., Jas, M., and Gramfort, A. (2018).
[Multivariate Convolutional Sparse Coding for Electromagnetic Brain Signals.](#)
In *Advances in Neural Information Processing Systems (NeurIPS)*, pages
3296–3306, Montreal, Canada

D-step: solving for the atoms

The dictionary update is performed by minimizing

$$\min_{\|\mathbf{D}_k\|_2 \leq 1} E(\{\mathbf{D}_k\}_k) \triangleq \sum_{n=1}^N \frac{1}{2} \|X^n - \sum_{k=1}^K z_k^n * \mathbf{D}_k\|_2^2 . \quad (1)$$

Computing $\nabla_{\mathbf{D}_k} E(\{\mathbf{D}_k\}_k)$ can be done efficiently

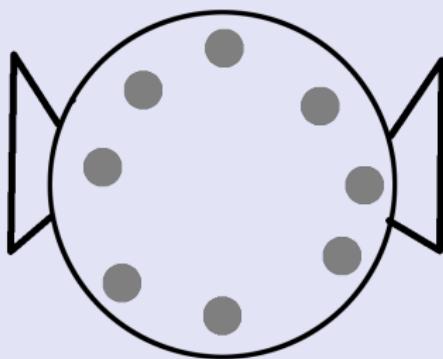
$$\nabla_{d_k} E(\{\mathbf{D}_k\}_k) = \sum_{n=1}^N (z_k^n)^\top * \left(X^n - \sum_{l=1}^K z_l^n * \mathbf{D}_l \right) = \Phi_k - \sum_{l=1}^K \Psi_{k,l} * \mathbf{D}_l ,$$

\Rightarrow Save with Projected Gradient Descent (PGD) with an Armijo backtracking line-search for the D-step [Wright and Nocedal, 1999].

However, this model does not account for the physics of the problem.


EM wave diffusion

- ▶ Recording here with 8 sensors




EM wave diffusion

- ▶ Recording here with 8 sensors
- ▶ EM activity in the brain



EM wave diffusion

- ▶ Recording here with 8 sensors
- ▶ EM activity in the brain
- ▶ The electric field is spread **linearly** and **instantaneously** over all sensors (Maxwell equations)



Multivariate CSC with rank-1 constraint

[Dupré la Tour et al. 2018, NeurIPS]

Idea: Impose a rank-1 constraint on the dictionary atoms D_k

To make the problem tractable, we decided to use auxiliary variables u_k and v_k s.t. $D_k = u_k v_k^\top$.

$$\begin{aligned} \min_{u_k, v_k, z_k^n} & \sum_{n=1}^N \frac{1}{2} \left\| X^n - \sum_{k=1}^K z_k^n * (u_k v_k^\top) \right\|_2^2 + \lambda \sum_{k=1}^K \|z_k^n\|_1, \\ \text{s.t. } & \|u_k\|_2^2 \leq 1, \|v_k\|_2^2 \leq 1 \text{ and } z_k^n \geq 0. \end{aligned} \tag{2}$$


Here,

- ▶ $u_k \in \mathbb{R}^P$ is the spatial pattern of our atom
- ▶ $v_k \in \mathbb{R}^L$ is the temporal pattern of our atom

⇒ Tri-convex optimization problem , solved with alternate minimization.

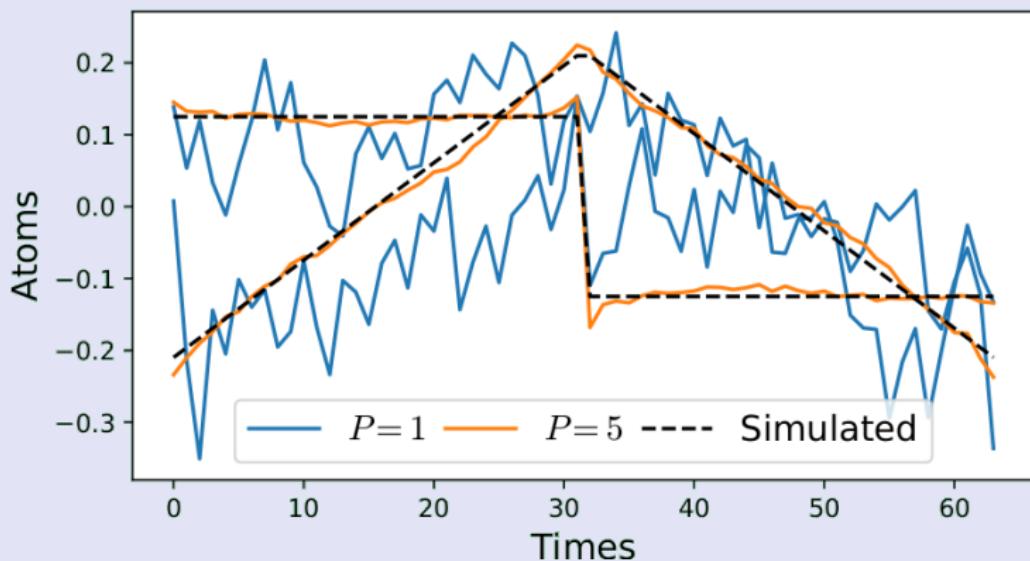
Fast optimization

Comparison with multivariate methods on somato dataset with
 $T = 134,700$, $K = 8$, $P = 5$ and $L = 128$



Pattern recovery


Patterns recovered with $P = 1$ and $P = 5$. The signals were generated with the two simulated temporal patterns and with $\sigma = 10^{-3}$.




MNE somatosensory data

A selection of temporal waveforms of the atoms learned on the MNE sample dataset.

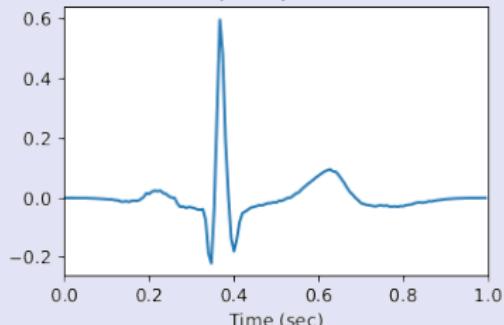
Spatial pattern 0
Explained variance 5.62 %




Spatial pattern 1
Explained variance 2.38 %



Temporal pattern 0



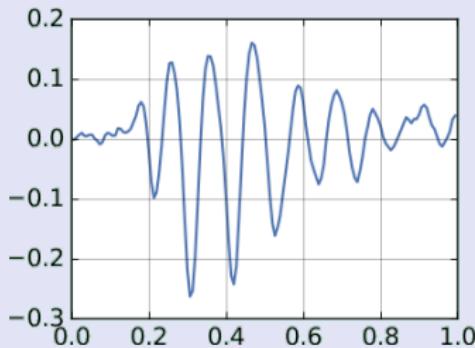


Temporal pattern 1



MNE somatosensory data

Atoms revealed using the MNE somatosensory data. Note the non-sinusoidal comb shape of the mu rhythm.



Recap Part III

Take home message

- ▶ The structure of the learned dictionary can be constrained to improve the interpretability of the recovered patterns.
- ▶ Can lead to more efficient algorithm and better recovery property.
- ▶ Open source package  <https://alphacsc.github.io>

Ahead of us

- ▶ Analysis of the patterns learned on large MEG database (HCP).
- ▶ Link between the learned waveforms and information propagation properties in the brain.
- ▶ Extension to scale invariant CDL to study frequency coupling in the brain.

Convolutional Dictionary Learning

- ▶ Flexible pattern extraction technique,
- ▶ Computationally tractable for more and more problems,
- ▶ Some application are already beginning to emerge.

Challenges

- ▶ Theoretical challenges remains (convergence, recoverability),
- ▶ The evaluation (and thus the parameter choices) is still not clear,
- ▶ Can give some insight for deep learning models?

Thanks!

Code available online:

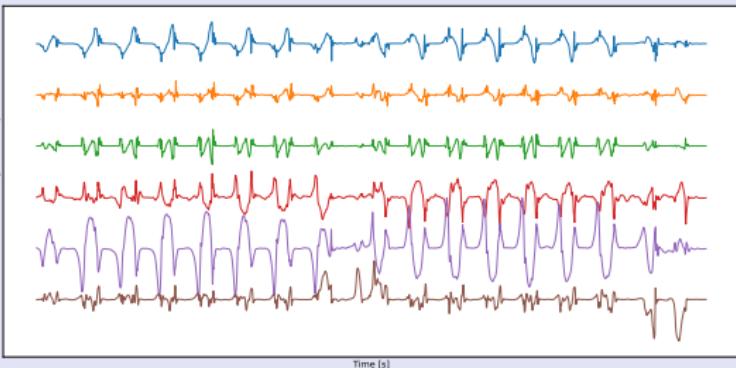
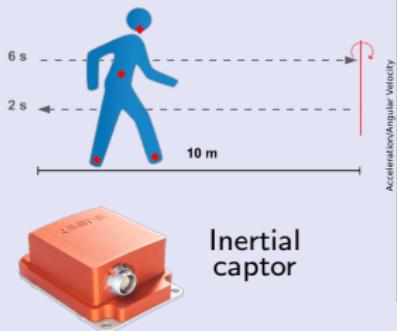
- ⌚ **LISTA** : github.com/tommoral/AdaptiveOptim
- ⌚ **DICOD** (& DiCoDiLe soon) : github.com/tommoral/dicod
- ⌚ **alphacsc** : alphacsc.github.io

Slides are on my web page:

 tommoral.github.io

 [@tomamoral](https://twitter.com/tomamoral)

Signals from human walking



- ▶ Shift invariant patterns linked to steps,
- ▶ Manual segmentation of the signal is expensive.

⇒ Can we do better with data-driven approach?

Experiment

Create a dictionary with 25 Gaussian patterns ($W = 90$)

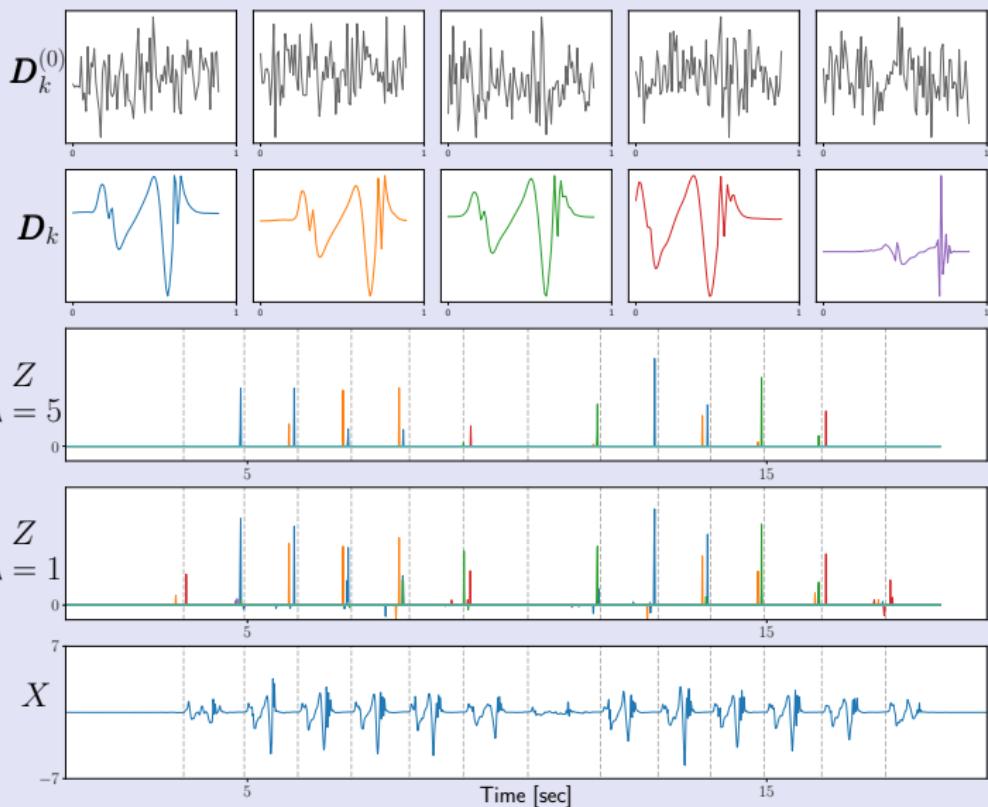
$$\mathcal{D}_k^{(0)} \sim \mathcal{N}(0, I_{90})$$

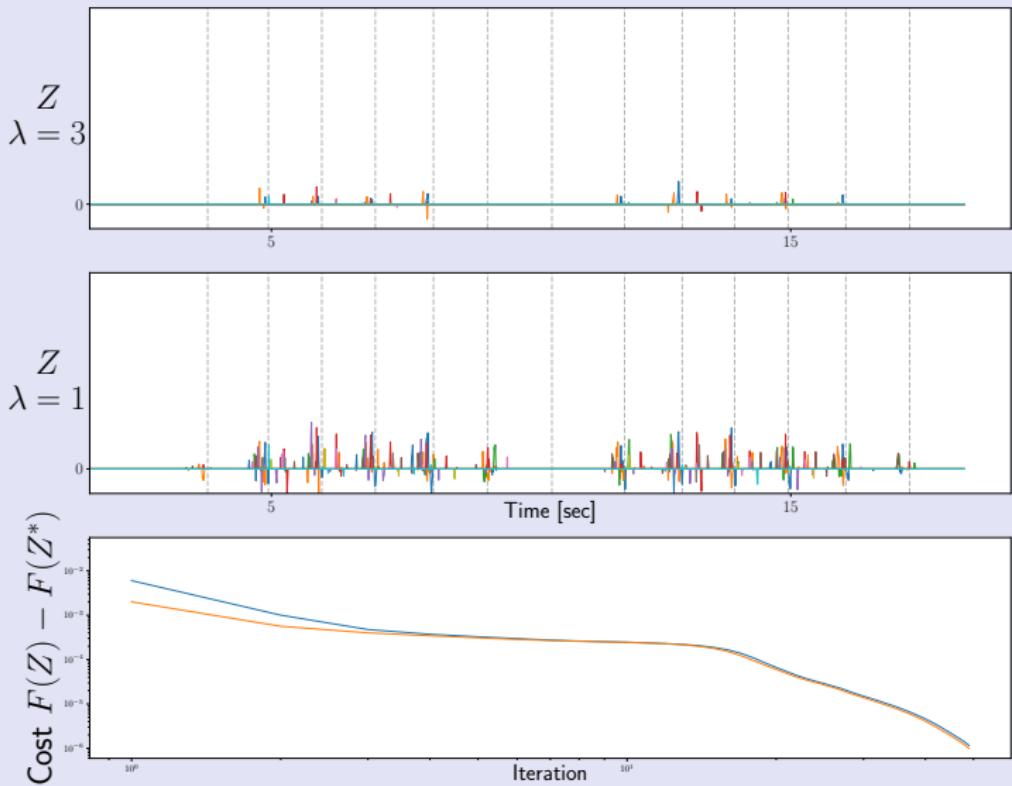
Use the Convolutional Dictionary Learning with DICOD to learn a dictionary \mathcal{D} on a set of 50 recording of healthy subjects walking.

Challenges

- ▶ Alignment of the patterns,
- ▶ Detect steps of different amplitude,
- ▶ Handle multivariate signals.

Experiment





Related works

- ▶ Giryes et al. [2018]: Propose the inexact projected gradient descent and conjecture that LISTA accelerate the LASSO resolution by learning the sparsity pattern of the input distribution.
- ▶ Xin et al. [2016]: Study the Hard-thresholding Algorithm and its capacity to recover the support of a sparse vector.
The paper relax the RIP conditions for the dictionary.

Generic Dictionaries

A dictionary $D \in \mathbb{R}^{p \times K}$ is a generic dictionary when its columns D_i are drawn uniformly over the ℓ_2 unit sphere \mathcal{S}^{p-1} .

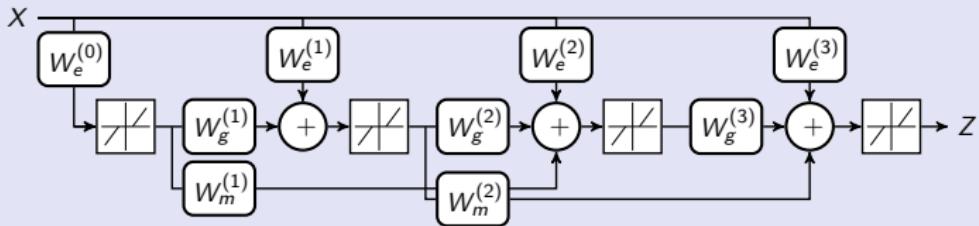
Theorem (Generic Acceleration)

In expectation over the generic dictionary D , the factorization algorithm using a diagonally dominant matrix $A \subset \mathcal{E}_\delta$, has better performance for iteration $q + 1$ than the normal ISTA iteration – which uses the identity – when

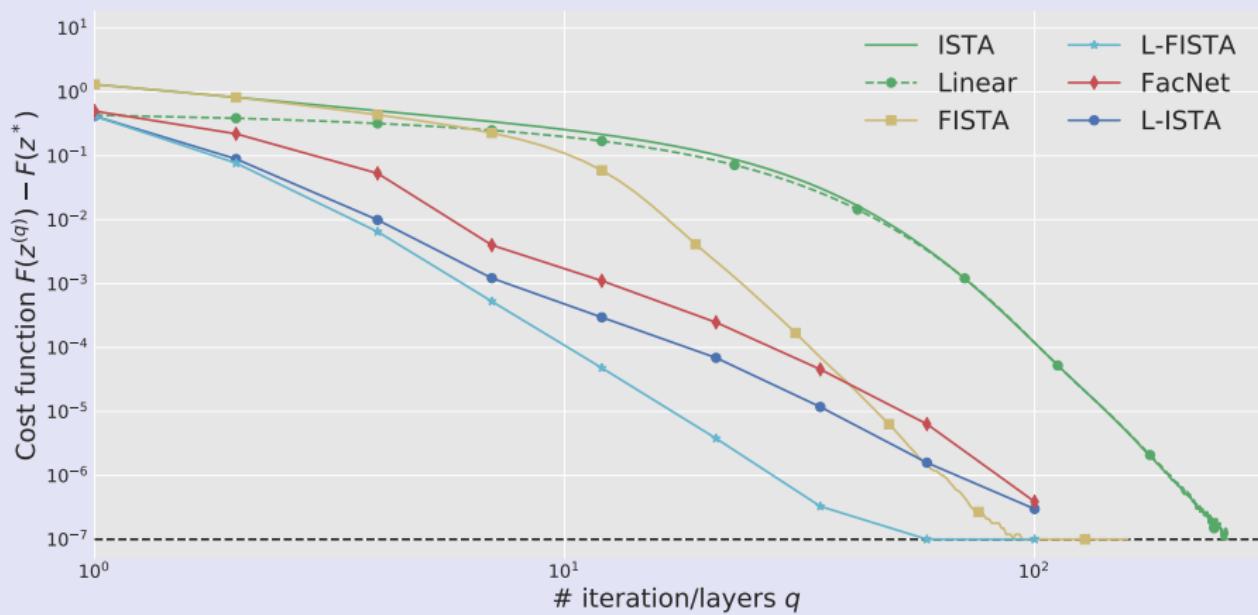
$$\lambda \mathbb{E}_z \left[\|z^{(q+1)}\|_1 + \|z^*\|_1 \right] \leq \sqrt{\frac{K(K-1)}{p}} \underbrace{\mathbb{E}_z \left[\|z^{(q)} - z^*\|_2^2 \right]}_{\text{expected resolution at iteration } q}$$

FacNet can improve the performances compared to ISTA when this is verified.

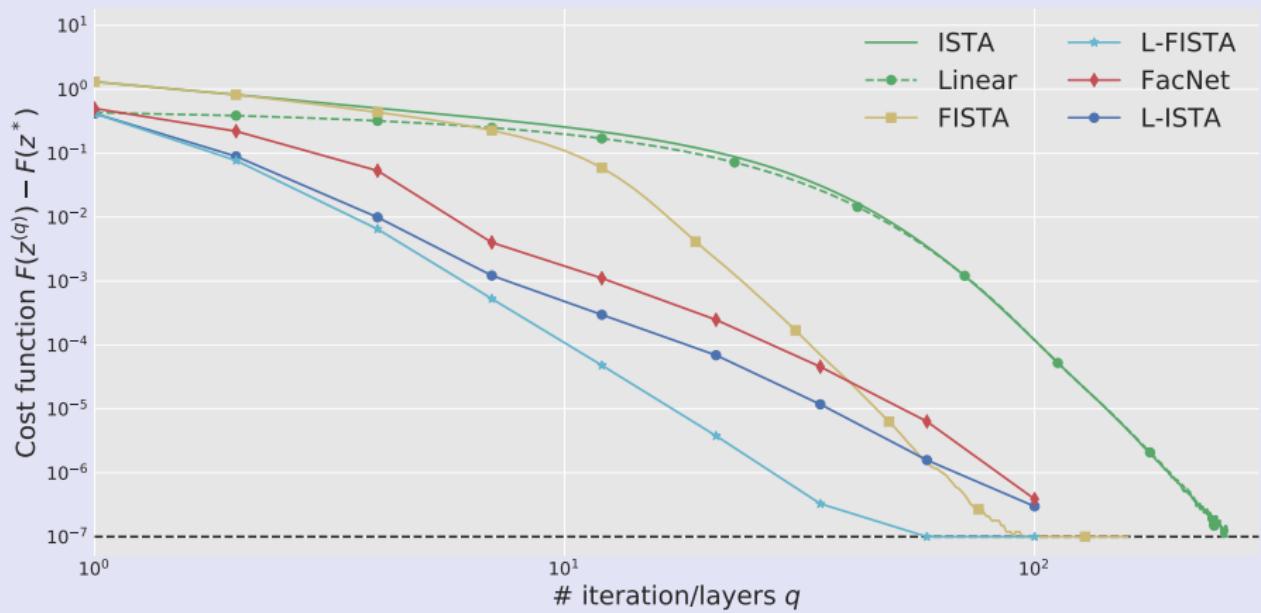
L-FISTA



Network architecture for L-FISTA.




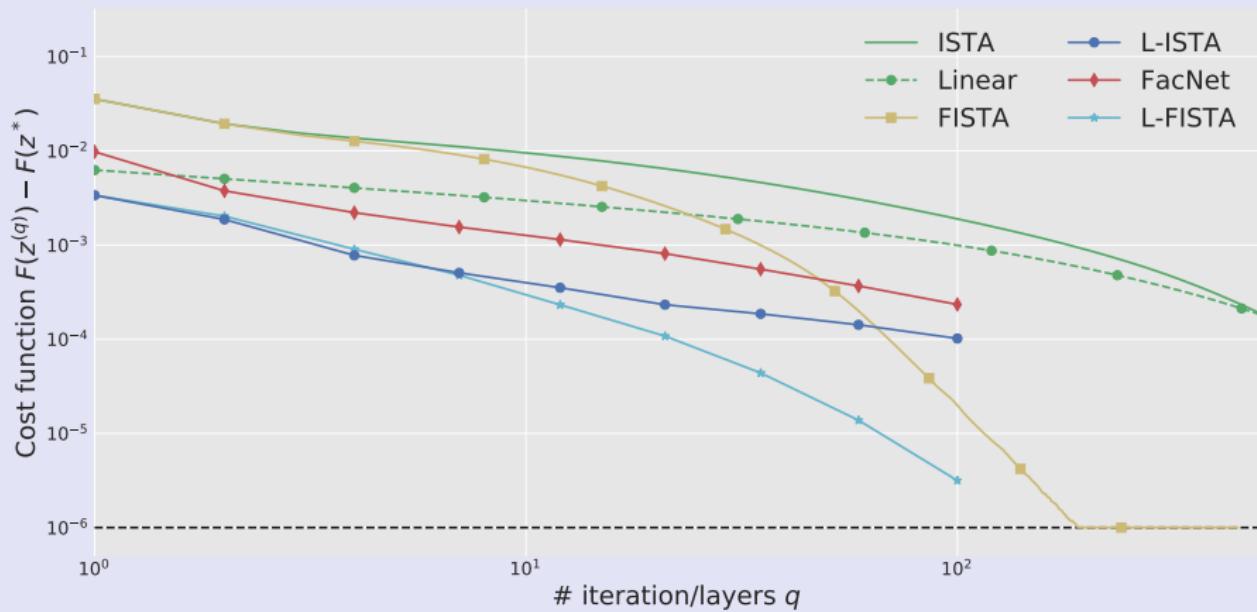
Evolution of the cost function $F(z^{(q)}) - F(z^*)$ with the number of layers/iterations q with a denser model



$$\rho = 1/20.$$

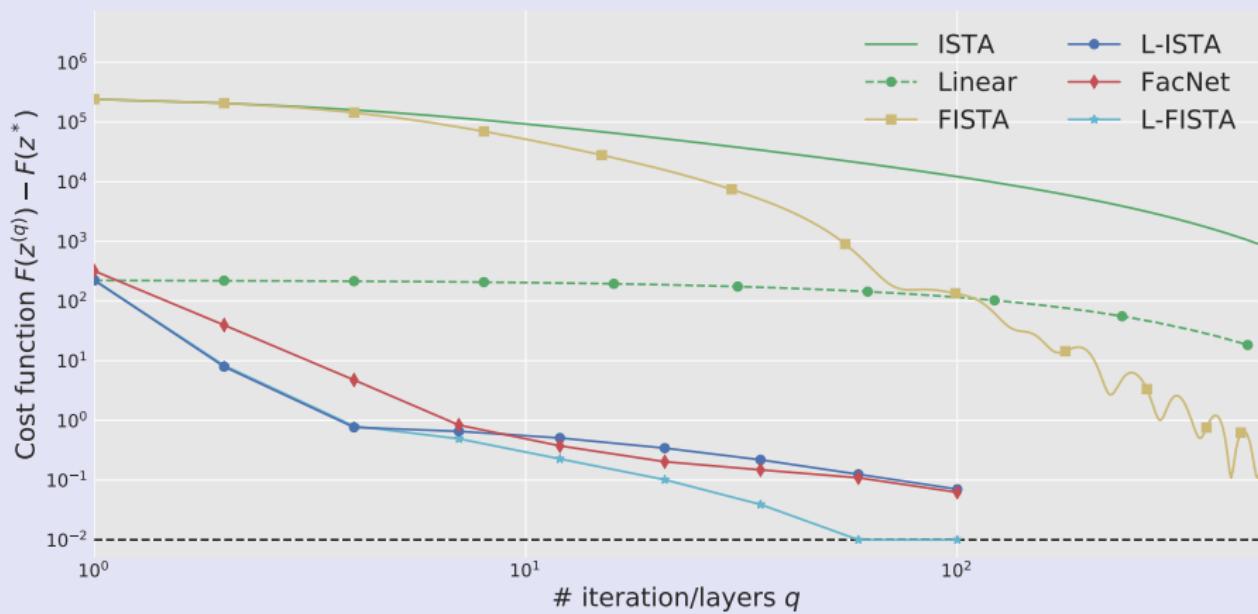
Evolution of the cost function $F(z^{(q)}) - F(z^*)$ with the number of layers/iterations q with a denser model





Evolution of the cost function $F(z^{(q)}) - F(z^*)$ with the number of layers

Dictionary D with $K = 100$ atoms learned on 10 000 MNIST samples (17x17) with dictionary learning. LISTA trained with MNIST training set and tested on MNIST test set.



Adversarial dictionary

The dictionary is constructed such that its eigen-vectors are sampled from the Fourier basis, with

$$D_{k,j} = e^{-2i\pi k \zeta_j}$$

for a random subset of frequencies

$$\{\zeta_i\}_{0 \leq i \leq p} \sim \mathcal{U}\left\{\frac{m}{K}; 0 \leq m \leq \frac{K}{2}\right\}$$

Diagonalizing B implies large deformation of the ℓ_1 -norm.

Finishing the process in a distributed environment

Non trivial point: **How to decide that the algorithm has converged?**

- ▶ Neighbors paused is not enough!
- ▶ Define a master 0 and send probes.
Wait for M probes return.
- ▶ Uses the notion of message queue and network flow.
Maybe we can have better way?

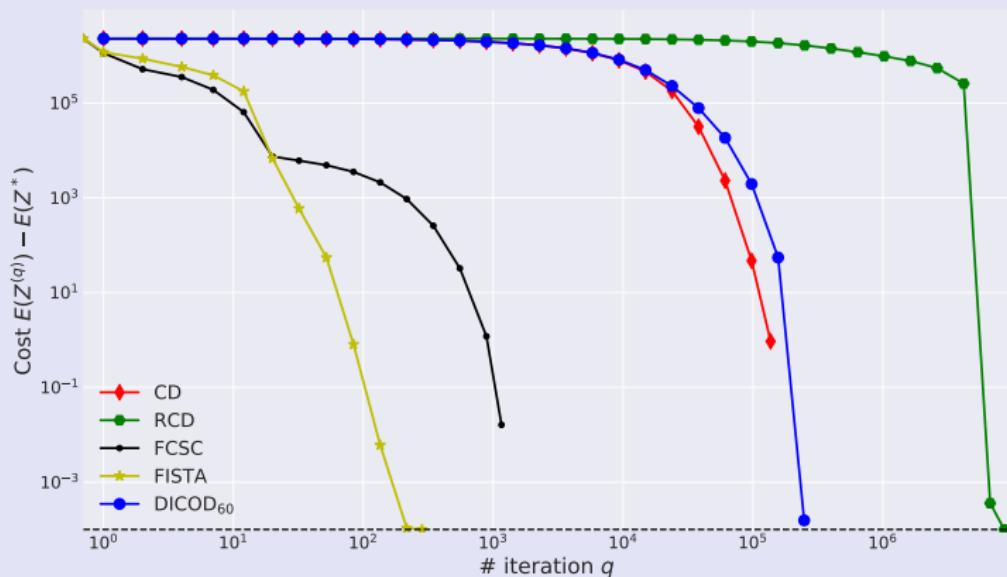
Numerical Experiments

Test on long signals generated with Bernoulli-Gaussian coding signal Z and a Gaussian dictionary \mathcal{D} . Fixed $K = 25$, $W = 200$ and $T = 600 * W$,

Algorithms implemented for benchmark

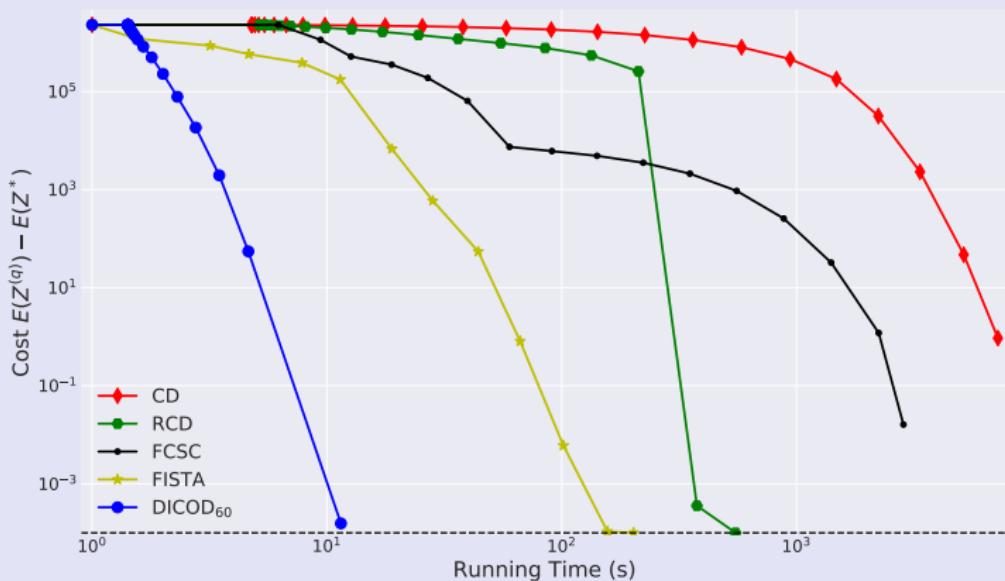
- ▶ Coordinate Descent (CD) [Kavukcuoglu et al., 2010]
- ▶ Randomized Coordinate Descent (RCD) [Nesterov, 2010]
- ▶ Fast Convolutional Sparse Coding (FCSC) [Bristow et al., 2013]
- ▶ Fast Iterative Soft-Thresholding Algorithm (FISTA) [Chalasani et al., 2013; Wohlberg, 2016]
- ▶ DICOD with 60 cores

Numerical convergence



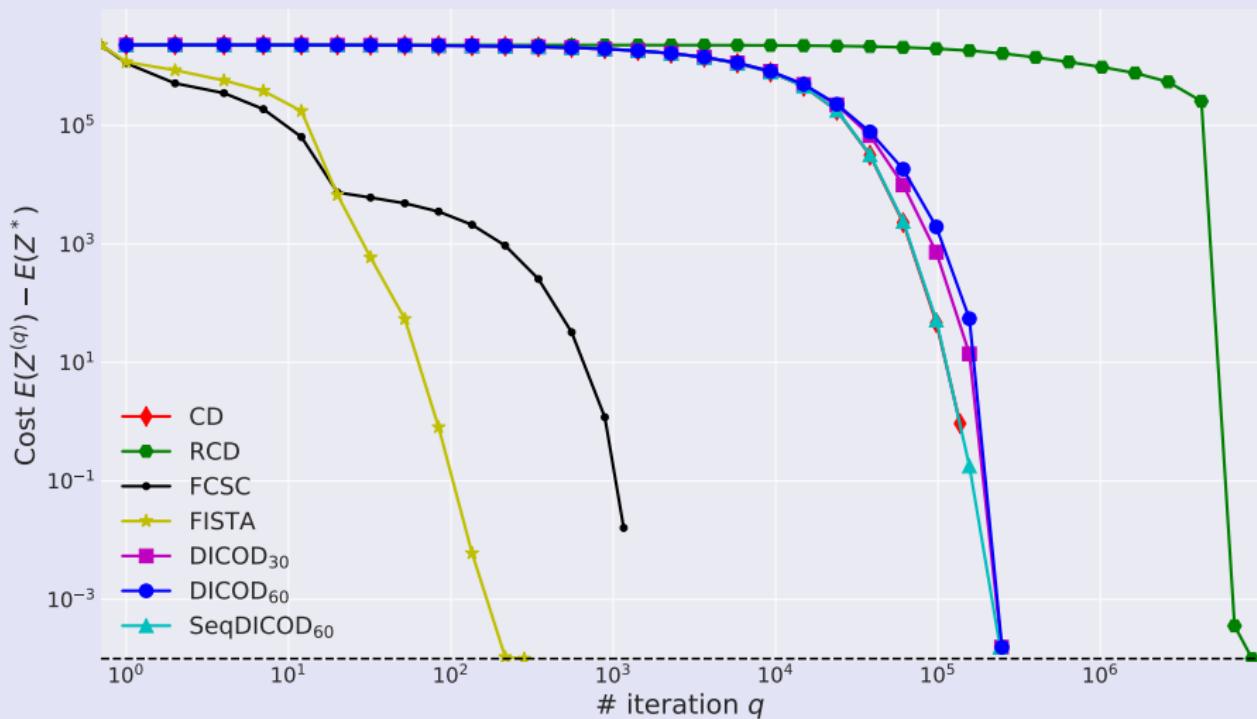
Cost as a function of the iterations

Numerical convergence



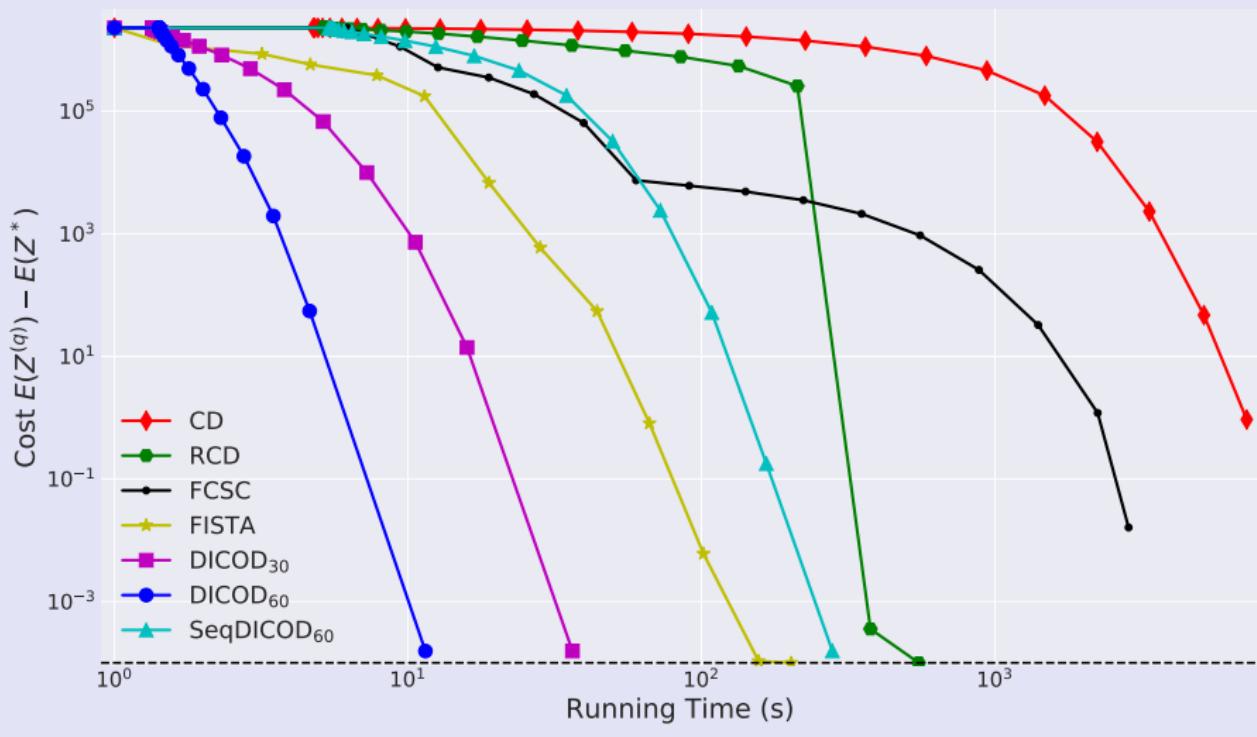
Cost as a function of the runtime

DICOD: numerical convergence



Cost as a function of the iterations

DICOD: numerical convergence



Cost as a function of the time

Complexity Analysis

Two sources of acceleration:

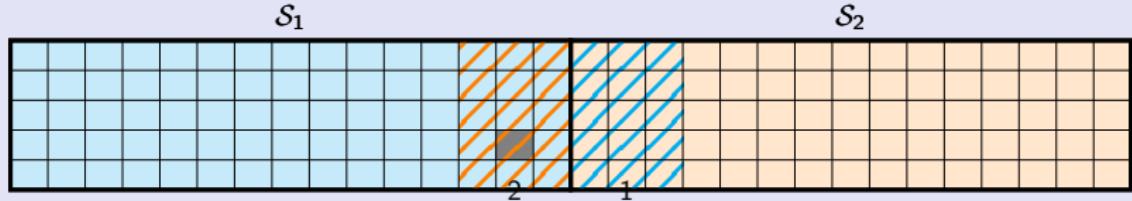
- ▶ Perform M updates in parallel,
- ▶ Each update is computed on a segment of size $\frac{L}{M}$
Iteration complexity of $\mathcal{O}\left(K \frac{L}{M}\right)$ instead of $\mathcal{O}(KL)$

Limitations:

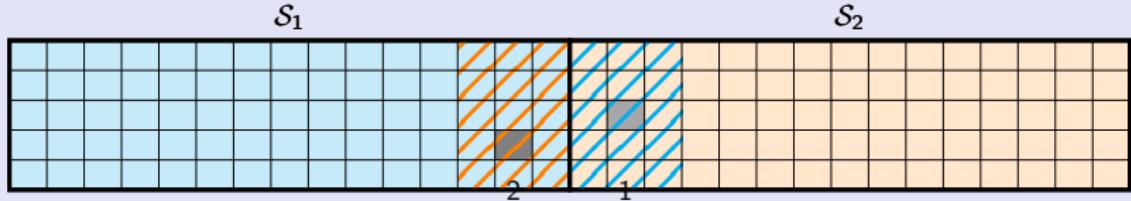
- ▶ Interfering updates, with probability $\alpha^2 = \left(\frac{WM}{T}\right)^2$
$$\mathbb{E}[Q_{dicod}] \underset{\alpha \rightarrow 0}{\gtrsim} M(1 - 2\alpha^2 M^2 + \mathcal{O}(\alpha^4 M^4)) .$$
- ▶ Cost of the update of β in $\mathcal{O}(KW)$



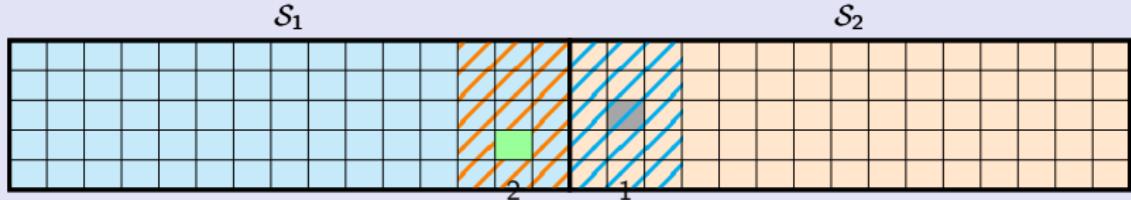
- ▶ Keep track of the value of the optimal update in an extended zone of size $L - 1$.



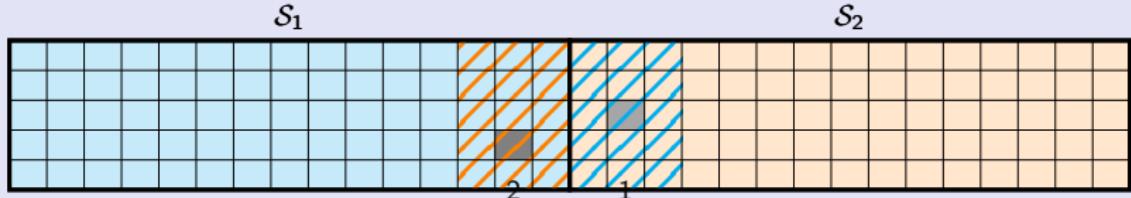
- ▶ Keep track of the value of the optimal update in an extended zone of size $L - 1$.
- ▶ Select an update candidate with LGCD.



- ▶ Keep track of the value of the optimal update in an extended zone of size $L - 1$.
- ▶ Select an update candidate with LGCD.
- ▶ If it is in the interfering zone, compare the value of the update with the value potential updates in the other worker.

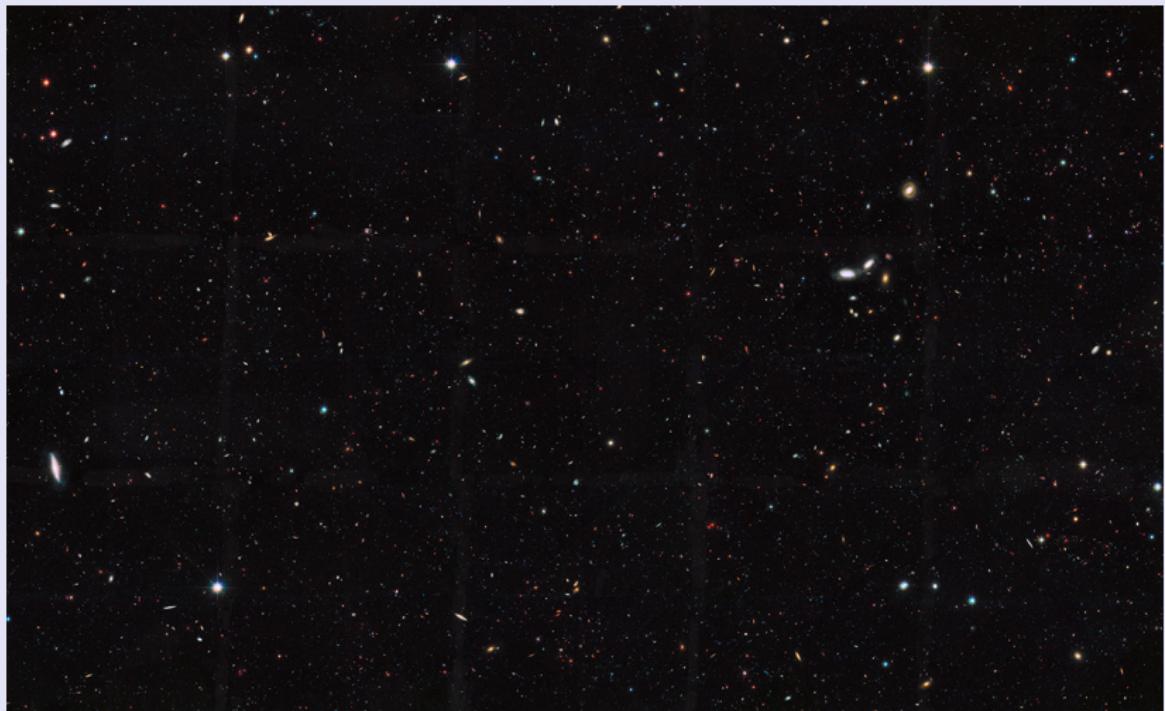


- ▶ Keep track of the value of the optimal update in an extended zone of size $L - 1$.
- ▶ Select an update candidate with LGCD.
- ▶ If it is in the interfering zone, compare the value of the update with the value potential updates in the other worker.
- ▶ Only perform the update if it is larger than the other update.




- ▶ Keep track of the value of the optimal update in an extended zone of size $L - 1$.
- ▶ Select an update candidate with LGCD.
- ▶ If it is in the interfering zone, compare the value of the update with the value potential updates in the other worker.
- ▶ Only perform the update if it is larger than the other update.
⇒ Give an update order asynchronously.

Images from Hubble Space Telescope

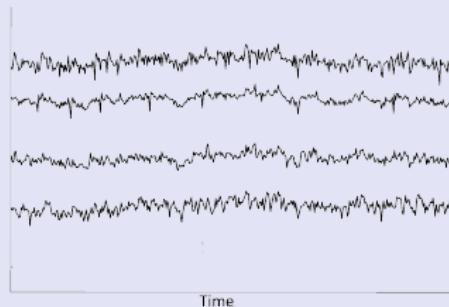



Images from Hubble Space Telescope



Studying brain activity through electromagnetic signals

- ▶ Brain (electrical) activity produces an electromagnetic field.
- ▶ This can be measured with EEG or MEG.



Goal: Study Oscillation in Neural Data

Oscillations are believed to play an important role in cognitive functions.

Many studies rely on Fourier or wavelet analyses:

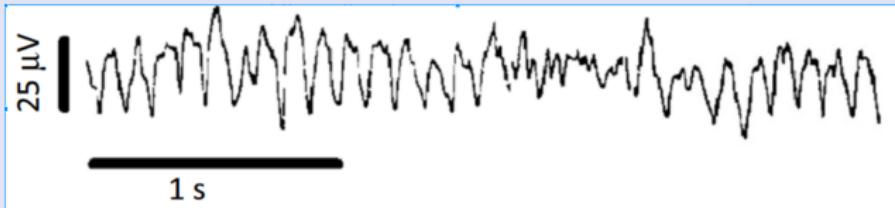
- ▶ Easy interpretation,
- ▶ Standard analysis e.g. canonical bands alpha, beta or theta.

[Buzsaki, 2006]


Goal: Study Oscillation in Neural Data

However, some brain rhythms are not sinusoidal, e.g. mu-waves.

[Hari, 2006]



and filtering degrades waveforms



⇒ Can we do better with data-driven approach?

Pattern recovery

Evolution of the recovery loss with σ for different values of P . Using more channels improves the recovery of the original patterns.

